

Questions

by Tong Wu, Lu Yu

1. Which capabilities API (seccomp-bpf, AppArmor, or SELinux) did you choose? Why did you make that choice?

We choose to use seccomp-bpf as our capabilities API.

Seccomp is a Linux feature that allows a userspace program to set up syscall filters. The BPF program loaded into the kernel starts with a system call + arguments and results in a filtering decision. Based on the results of the filter, the system call can be allowed, blocked, or the process can be killed. It can be easily implemented with the `<seccomp-bpf.h>` header file included in the program code.

AppArmor, or SELinux is MAC (Mandatory Access Control) systems/framework in which you add policies to sandboxing. It is used to whitelist or blacklist a subject's (program's) access to an object (file, path, etc.). The policies can also be used to restrict capabilities, or even limit network access. These strive for system-wide enforcement of policies that control the actions and resources that each program on a system can perform. Comparing to AppArmor, SELinux can be more complex to configure.

In this assignment, our goal is to restrict system calls in the program. There is no need to add a systemwide security policy. With seccomp-bpf, the requirements can be achieved by just modifying the client and server code. However, if using AppArmor or SELinux, a separate configuration file needs to be added, and a sandbox or VM is required. Otherwise, the policy will be implemented in the entire system. This costs much more effort than seccomp-bpf.

2. What was the process you used to ascertain the list of system calls required by each program?

- a. use `strace -x fc ./a.out` to get the required system calls for each program
- b. `#include <seccomp.h>`
- c. Initiate the filter with `scmp_filter_ctx ctx` and
`ctx = seccomp_init(SCMP_ACT_TRAP)`
- d. Add the list of allowed system calls to the filter with `seccomp_rule_add()`
- e. Load the filter with `seccomp_load(ctx)`

3.What system calls are needed by each?

client:

% time	seconds	usecs/call	calls	errors	syscall
0.00	0.000000	0	3		read
0.00	0.000000	0	4		write
0.00	0.000000	0	2		close
0.00	0.000000	0	3		fstat
0.00	0.000000	0	5		mmap
0.00	0.000000	0	4		mprotect
0.00	0.000000	0	1		munmap
0.00	0.000000	0	3		brk
0.00	0.000000	0	3	3	access
0.00	0.000000	0	1		socket
0.00	0.000000	0	1		connect
0.00	0.000000	0	1		sendto
0.00	0.000000	0	1		execve
0.00	0.000000	0	1		arch_prctl
0.00	0.000000	0	2		openat
100.00	0.000000		35	3	total

server:

% time	seconds	usecs/call	calls	errors	syscall
20.56	0.000022	1	23		mmap
19.63	0.000021	1	16		mprotect
11.21	0.000012	6	2	1	wait4
9.35	0.000010	1	10	10	access
8.41	0.000009	1	12		openat
7.48	0.000008	3	3		socket
5.61	0.000006	1	5		munmap
3.74	0.000004	0	14		close
3.74	0.000004	0	13		fstat
2.80	0.000003	0	9		read
1.87	0.000002	1	4		brk
1.87	0.000002	2	1		clone
1.87	0.000002	1	2		execve
1.87	0.000002	1	2		arch_prctl
0.00	0.000000	0	6		write
0.00	0.000000	0	1		stat
0.00	0.000000	0	2		lseek
0.00	0.000000	0	2	2	connect
0.00	0.000000	0	1		accept
0.00	0.000000	0	1		sendto
0.00	0.000000	0	1		bind
0.00	0.000000	0	1		listen
0.00	0.000000	0	1		setsockopt
0.00	0.000000	0	1		chdir
0.00	0.000000	0	2		getuid
0.00	0.000000	0	1		setuid
0.00	0.000000	0	1		chroot
100.00	0.000107		137	13	total

4.What happens when your application calls the prohibited system call? What is the application behaviour that results from the call?

In strict mode, seccomp kills a process when it violates the policy. However, seccomp-bpf allows a number of actions to take based on the results of running the policy:

- Killing the process
- Sending the process a SIGSYS signal
- Failing the system call and returning a (filter-provided) errno value
- Notifying an attached process tracer (ptrace()), if one is attached. In turn, the process tracer can skip or even change the system call.
- Allowing the system call